

IN3050/IN4050 Mandatory Assignment 3: Unsupervised Learning

Name: Elias Goldmann Brandtzæg

Username: eliasgb

Rules

Before you begin the exercise, review the rules at this website: <https://www.uio.no/english/studies/examinations/compulsory-activities/mn-ifi-mandatory.html> , in particular the paragraph on cooperation. This is an individual assignment. You are not allowed to deliver together or copy/share source-code/answers with others. Read also the "Routines for handling suspicion of cheating and attempted cheating at the University of Oslo" <https://www.uio.no/english/about/regulations/studies/studies-examinations/routines-cheating.html> By submitting this assignment, you confirm that you are familiar with the rules and the consequences of breaking them.

Delivery

Deadline: Friday, April 21, 2023, 23:59

Your submission should be delivered in Devilry. You may redeliver in Devilry before the deadline, but include all files in the last delivery, as only the last delivery will be read. You are recommended to upload preliminary versions hours (or days) before the final deadline.

What to deliver?

You are recommended to solve the exercise in a Jupyter notebook, but you might solve it in a Python program if you prefer.

If you choose Jupyter, you should deliver the notebook. You should answer all questions and explain what you are doing in Markdown. Still, the code should be properly commented. The notebook should contain results of your runs. In addition, you should make a pdf of your solution which shows the results of the runs.

If you prefer not to use notebooks, you should deliver the code, your run results, and a pdf-report where you answer all the questions and explain your work.

Your report/notebook should contain your name and username.

Deliver one single zipped folder (.zip, .tgz or .tar.gz) which contains your complete solution.

Important: if you weren't able to finish the assignment, use the PDF report/Markdown to elaborate on what you've tried and what problems you encountered. Students who have made an effort and attempted all parts of the assignment will get a second chance even if they fail initially. This exercise will be graded PASS/FAIL.

Goals of the exercise

This exercise has three parts. The first part is focused on Principal Component Analysis (PCA). You will go through some basic theory, and implement PCA from scratch to do compression and visualization of data.

The second part focuses on clustering using K-means. You will use `scikit-learn` to run K-means clustering, and use PCA to visualize the results.

The last part ties supervised and unsupervised learning together in an effort to evaluate the output of K-means using a logistic regression for multi-class classification approach.

The master students will also have to do one extra part about tuning PCA to balance compression with information lost.

Tools

You may freely use code from the weekly exercises and the published solutions. In the first part about PCA you may **NOT** use ML libraries like `scikit-learn`. In the K-means part and beyond we encourage the use of `scikit-learn` to iterate quickly on the problems.

Beware

This is a new assignment. There might occur typos or ambiguities. If anything is unclear, do not hesitate to ask. Also, if you think some assumptions are missing, make your own and explain them!

Principal Component Analysis (PCA)

In this section, you will work with the PCA algorithm in order to understand its definition and explore its uses.

Implementation: how is PCA implemented?

Here we implement the basic steps of PCA and we assemble them.

Importing libraries

We start importing the *numpy* library for performing matrix computations, the *pyplot* library for plotting data, and the *syntheticdata* module to import synthetic data.

```
In [324... import numpy as np
import matplotlib.pyplot as plt

import syntheticdata
```

Centering the Data

Implement a function with the following signature to center the data as explained in *Marsland*.

```
In [325... def center_data(A):
    # INPUT:
    # A    [NxM] numpy data matrix (N samples, M features)
    #
    # OUTPUT:
    # X    [NxM] numpy centered data matrix (N samples, M features)

    # Computing the mean value of each column
    mean = np.mean(A, axis=0)
    data = A.copy()
    # Centering the data by subtracting from each column with the mean value of all the values in the column
```

```
data -= mean
return data
```

Test your function checking the following assertion on *testcase*:

```
In [326... testcase = np.array([[3., 11., 4.3], [4., 5., 4.3], [5., 17., 4.5], [4, 13., 4.4]])
answer = np.array([[-1., -0.5, -0.075], [0., -6.5, -0.075], [1., 5.5, 0.125], [0., 1.5, 0.025]])
np.testing.assert_array_almost_equal(center_data(testcase), answer)
```

Computing Covariance Matrix

Implement a function with the following signature to compute the covariance matrix as explained in *Marsland*.

```
In [327... def compute_covariance_matrix(A):
    # INPUT:
    # A      [NxM] centered numpy data matrix (N samples, M features)
    #
    # OUTPUT:
    # C      [MxM] numpy covariance matrix (M features, M features)
    #
    # Do not apply centering here. We assume that A is centered before this function is called.
    data = A
    C = np.cov(np.transpose(data))

    return C
```

Test your function checking the following assertion on *testcase*:

```
In [328... testcase = center_data(np.array([[22., 11., 5.5], [10., 5., 2.5], [34., 17., 8.5], [28., 14., 7]]))
answer = np.array([[580., 290., 145.], [290., 145., 72.5], [145., 72.5, 36.25]])

# Depending on implementation the scale can be different:
to_test = compute_covariance_matrix(testcase)

answer = answer/answer[0, 0]
to_test = to_test/to_test[0, 0]
```

```
np.testing.assert_array_almost_equal(to_test, answer)
```

Computing eigenvalues and eigenvectors

Use the linear algebra package of `numpy` and its function `np.linalg.eig()` to compute eigenvalues and eigenvectors. Notice that we take the real part of the eigenvectors and eigenvalues. The covariance matrix *should* be a symmetric matrix, but the actual implementation in `compute_covariance_matrix()` can lead to small round off errors that lead to tiny imaginary additions to the eigenvalues and eigenvectors. These are purely numerical artifacts that we can safely remove.

Note: If you decide to NOT use `np.linalg.eig()` you must make sure that the eigenvalues you compute are of unit length!

```
In [329... def compute_eigenvalue_eigenvectors(A):
# INPUT:
# A      [DxD] numpy matrix
#
# OUTPUT:
# eigval  [D] numpy vector of eigenvalues
# eigvec  [DxD] numpy array of eigenvectors

# Computing eigenvalues and eigenvectors
data = A
eigval, eigvec = np.linalg.eig(data)

# Numerical roundoff can lead to (tiny) imaginary parts. We correct that here.
eigval = eigval.real
eigvec = eigvec.real

return eigval, eigvec
```

Test your function checking the following assertion on *testcase*:

```
In [330... testcase = np.array([[2,0,0],[0,5,0],[0,0,3]])
answer1 = np.array([2.,5.,3.])
answer2 = np.array([[1.,0.,0.],[0.,1.,0.],[0.,0.,1.]])
x,y = compute_eigenvalue_eigenvectors(testcase)
```

```
np.testing.assert_array_almost_equal(x, answer1)
np.testing.assert_array_almost_equal(y, answer2)
```

Sorting eigenvalues and eigenvectors

Implement a function with the following signature to sort eigenvalues and eigenvectors as explained in *Marsland*.

Remember that eigenvalue $eigval[i]$ corresponds to eigenvector $eigvec[:,i]$.

```
In [331... def sort_eigenvalue_eigenvectors(eigval, eigvec):
    # INPUT:
    # eigval      [D] numpy vector of eigenvalues
    # eigvec      [DxD] numpy array of eigenvectors
    #
    # OUTPUT:
    # sorted_eigval      [D] numpy vector of eigenvalues
    # sorted_eigvec      [DxD] numpy array of eigenvectors

    indices = np.argsort(eigval)
    indices = indices[::-1]

    sorted_eigval = eigval[indices]
    sorted_eigvec = eigvec[:,indices]

    return sorted_eigval, sorted_eigvec
```

Test your function checking the following assertion on *testcase*:

```
In [332... testcase = np.array([[2,0,0],[0,5,0],[0,0,3]])
answer1 = np.array([5.,3.,2.])
answer2 = np.array([[0.,0.,1.],[1.,0.,0.],[0.,1.,0.]])
x,y = compute_eigenvalue_eigenvectors(testcase)
x,y = sort_eigenvalue_eigenvectors(x,y)
np.testing.assert_array_almost_equal(x, answer1)
np.testing.assert_array_almost_equal(y, answer2)
```

PCA Algorithm

Implement a function with the following signature to compute PCA as explained in *Marsland* using the functions implemented above.

```
In [333... def pca(A,m):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    # m      integer number denoting the number of learned features (m <= M)
    #
    # OUTPUT:
    # pca_eigvec      [Mxm] numpy matrix containing the eigenvectors (M dimensions, m eigenvectors)
    # P               [Nxm] numpy PCA data matrix (N samples, m features)

    # Centering the data

    data = center_data(A)

    # Computing a covariance matrix
    C = compute_covariance_matrix(data)

    # Computing eigenvalues and eigenvectors
    eigval, eigvec = compute_eigenvalue_eigenvectors(C)

    # Sorting the eigenvalues and eigenvectors
    sorted_eigval, sorted_eigvec = sort_eigenvalue_eigenvectors(eigval, eigvec)

    # Reducing the dimensions by m
    if m > 0:
        sorted_eigvec = sorted_eigvec[:, :m]

    pca_eigvec = sorted_eigvec
    P = np.dot(np.transpose(sorted_eigvec), np.transpose(data))

    return pca_eigvec, P.T
```

Test your function checking the following assertion on *testcase*:

```
In [334... testcase = np.array([[22.,11.,5.5],[10.,5.,2.5],[34.,17.,8.5]])
x,y = pca(testcase,2)

import pickle
answer1_file = open('PCAanswer1.pkl','rb'); answer2_file = open('PCAanswer2.pkl','rb')
answer1 = pickle.load(answer1_file); answer2 = pickle.load(answer2_file)

test_arr_x = np.sum(np.abs(np.abs(x) - np.abs(answer1)), axis=0)
np.testing.assert_array_almost_equal(test_arr_x, np.zeros(2))

test_arr_y = np.sum(np.abs(np.abs(y) - np.abs(answer2)))
np.testing.assert_almost_equal(test_arr_y, 0)
```

Understanding: how does PCA work?

We now use the PCA algorithm you implemented on a toy data set in order to understand its inner workings.

Loading the data

The module *syntheticdata* provides a small synthetic dataset of dimension [100x2] (100 samples, 2 features).


```
In [335... X = syntheticdata.get_synthetic_data1()
```

Visualizing the data

Visualize the synthetic data using the function *scatter()* from the *matplotlib* library.

```
In [336... plt.scatter(X[:,0],X[:,1])
```

```
Out[336]: <matplotlib.collections.PathCollection at 0x7fc8645fc220>
```





Visualize the centered data

Notice that the data visualized above is not centered on the origin (0,0). Use the function defined above to center the data, and the replot it.

```
In [337... X_center = center_data(X)
plt.scatter(X_center[:,0], X_center[:,1])
```

```
Out[337]: <matplotlib.collections.PathCollection at 0x7fc843a08b50>
```



Visualize the first eigenvector

Visualize the vector defined by the first eigenvector. To do this you need:

- Use the `PCA()` function to recover the eigenvectors
- Plot the centered data as done above
- The first eigenvector is a 2D vector (x_0, y_0) . This defines a vector with origin in $(0,0)$ and head in (x_0, y_0) . Use the function `plot()` from matplotlib to plot a line over the first eigenvector.


```
In [338]: pca_eigvec, _ = pca(X_center, 0)
print("eigvec shape = {}".format(pca_eigvec.shape))
first_eigvec = pca_eigvec[:,0]

plt.scatter(X_center[:,0], X_center[:,1])

x = np.linspace(-5, 5, 1000)
y = first_eigvec[1]/first_eigvec[0] * x
plt.plot(x, y)
```

```
eigvec shape = (2, 2)
```

```
Out[338]: [<matplotlib.lines.Line2D at 0x7fc844ed1eb0>]
```




Visualize the PCA projection

Finally, use the `PCA()` algorithm to project on a single dimension and visualize the result using again the `scatter()` function.

```
In [339... x3 = X
_, P = pca(x3, 1)
plt.scatter(P, np.zeros_like(P))
```

```
Out[339]: <matplotlib.collections.PathCollection at 0x7fc8724e2040>
```



Evaluation: when are the results of PCA sensible?

So far we have used PCA on synthetic data. Let us now imagine we are using PCA as a pre-processing step before a classification task. This is a common setup with high-dimensional data. We explore when the use of PCA is sensible.

Loading the first set of labels

The function `get_synthetic_data_with_labels1()` from the module `syntethicdata` provides a first labeled dataset.

```
In [340... X,y = syntheticdata.get_synthetic_data_with_labels1()
```


Running PCA

Process the data using the PCA algorithm and project it in one dimension. Plot the labeled data using `scatter()` before and after running PCA. Comment on the results.

```
In [341... plt.scatter(X[:,0],X[:,1],c=y[:,0])  
plt.figure()
```

```
_, P = pca(X, 1)  
plt.scatter(P, np.ones(P.shape[0]), c=y[:, 0])
```

Out[341]: <matplotlib.collections.PathCollection at 0x7fc843a26b50>



Comment: We start with a dataset with two classes that are easy to separate. After running the PCA algorithm we can see the dataland scape for the two classes in one dimension. When seeing the data in two dimensions it's difficult to spot any overlapping data. In a one dimension we can observe that the data is indeed overlapping in some areas. In other words we could say that we found some underlying patterns when visualizing the data in one dimension.

Loading the second set of labels

The function `get_synthetic_data_with_labels2()` from the module `syntethicdata` provides a second labeled dataset.


```
In [342... X,y = syntheticdata.get_synthetic_data_with_labels2()
```


Running PCA

As before, process the data using the PCA algorithm and project it in one dimension. Plot the labeled data using `scatter()` before and after running PCA. Comment on the results.

```
In [343... plt.scatter(X[:,0],X[:,1],c=y[:,0])  
plt.figure()  
_,P = pca(X,1)  
plt.scatter(P,np.ones(P.shape[0]),c=y[:,0])
```

```
Out[343]: <matplotlib.collections.PathCollection at 0x7fc850ec0430>
```





Comment: When projecting the data in one dimension we experience more noise in the data visualization. It was easier to grasp the data before reducing the dimensions.

How would the result change if you were to consider the second eigenvector? Or if you were to consider both eigenvectors?

Answer: If we were to consider the second eigenvector we would get results of a different direction in the reduced-dimensional space, which may have captured additional patterns or structures in the data.

If we considered both eigenvectors we would observe the data in two dimensions. This would likely result in a more informative visualization of the data compared to one-dimensional projection.

Case study 1: PCA for visualization

We now consider the *iris* dataset, a simple collection of data ($N=150$) describing iris flowers with four ($M=4$) features. The features are: Sepal Length, Sepal Width, Petal Length and Petal Width. Each sample has a label, identifying each flower as one of 3 possible types of iris: Setosa, Versicolour, and Virginica.

Visualizing a 4-dimensional dataset is impossible; therefore we will use PCA to project our data in 2 dimensions and visualize it.

Loading the data

The function `get_iris_data()` from the module `syntethicdata` returns the *iris* dataset. It returns a data matrix of dimension `[150x4]` and a label vector of dimension `[150]`.


```
In [344... X,y = synteticdata.get_iris_data()
```

Visualizing the data by selecting features

Try to visualize the data (using label information) by randomly selecting two out of the four features of the data. You may try different pairs of features.

```
In [345... import random  
plt.scatter(X[:,1],X[:,2])
```

```
Out[345]: <matplotlib.collections.PathCollection at 0x7fc850ecc550>
```




Visualizing the data by PCA

Process the data using PCA and visualize it (using label information). Compare with the previous visualization and comment on the results.


```
In [346]: _, P = pca(X, 0)
plt.scatter(P[:,1], P[:,2])
```

```
Out[346]: <matplotlib.collections.PathCollection at 0x7fc86673bc70>
```



Comment: The previous visualization had several overlapping datapoints. It also seems like the dataset was distributed in diagonal lines. After using PCA the data got more scattered.

Case study 2: PCA for compression

We now consider the *faces in the wild (lfw)* dataset, a collection of pictures (N=1280) of people. Each pixel in the image is a feature (M=2914).

Loading the data

The function `get_lfw_data()` from the module `syntethicdata` returns the *lfw* dataset. It returns a data matrix of dimension [1280x2914] and a label vector of dimension [1280]. It also returns two parameters, h and w , reporting the height and the width of the images (these parameters are necessary to plot the data samples as images). Beware, it might take some time to download the data. Be patient :)

```
In [347]: X,y,h,w = syntethicdata.get_lfw_data()
```

Inspecting the data

Choose one datapoint to visualize (first coordinate of the matrix X) and use the function `imshow()` to plot and inspect some of the pictures.

Notice that `imshow` receives as a first argument an image to be plot; the image must be provided as a rectangular matrix, therefore we reshape a sample from the matrix X to have height h and width w . The parameter `cmap` specifies the color coding; in our case we will visualize the image in black-and-white with different gradations of grey.

```
In [348... plt.imshow(X[1,:].reshape((h, w)), cmap=plt.cm.gray)
```

```
Out[348]: <matplotlib.image.AxesImage at 0x7fc8508efaf0>
```



Implementing a compression-decompression function

Implement a function that first uses PCA to project samples in low-dimensions, and then reconstruct the original image.

Hint: Most of the code is the same as the previous `PCA()` function you implemented. You may want to refer to *Marsland* to check out how reconstruction is performed.

```
In [349... def encode_decode_pca(A,m):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    # m      integer number denoting the number of learned features (m <= M)
    #
    # OUTPUT:
    # Ahat   [NxM] numpy PCA reconstructed data matrix (N samples, M features)

    data = A.copy()
    # Perform pca projection on data set
    sorted_eigvec, P = pca(data, m)
    # Inverse the transformation to get back to original
    #Ahat = np.dot(P.T, data) + np.mean(data, axis=0)
    # Reconstruct the original data matrix from the encoded data matrix
    Ahat = np.dot(P, sorted_eigvec.T) + np.mean(data, axis=0)

    #np.dot(np.transpose(sorted_eigvec),np.transpose(data))

    return Ahat
```

Compressing and decompressing the data

Use the implemented function to encode and decode the data by projecting on a lower dimensional space of dimension 200 (m=200).

```
In [350... Xhat = encode_decode_pca(X,200)
```

Inspecting the reconstructed data


Use the function *imshow* to plot and compare original and reconstructed pictures. Comment on the results.

```
In [351... fig, axs = plt.subplots(1, 2, figsize=(10, 5))
print("X shape = {}".format(X.shape))
print("Xhat shape = {}".format(Xhat.shape))
axs[0].imshow(X[2,:].reshape((h, w)), cmap=plt.cm.gray)
axs[0].set_title("Original image")
axs[1].imshow(Xhat[2,:].reshape((h, w)), cmap=plt.cm.gray)
axs[1].set_title("Reconstructed image")
```

```
X shape = (1288, 2914)
```

```
Xhat shape = (1288, 2914)
```

```
Out[351]: Text(0.5, 1.0, 'Reconstructed image')
```



Comment: The reconstructed image looks similar to the original image. We can see that the quality is not as good and the cause could be that we've lost some important features during the compression of the image.


Evaluating different compressions

Use the previous setup to generate compressed images using different values of low dimensions in the PCA algorithm (e.g.: 100, 200, 500, 1000). Plot and comment on the results.

```
In [352]: plt.imshow(X[0,:].reshape((h, w)), cmap=plt.cm.gray)

plt.figure()
Xhat = encode_decode_pca(X,1000)
plt.imshow(Xhat[0,:].reshape((h, w)), cmap=plt.cm.gray)
```

```
Out[352]: <matplotlib.image.AxesImage at 0x7fc8506e0850>
```



Comment: When the value `m` is increasing, the quality of the reconstructed image is getting better.

Master Students: PCA Tuning

If we use PCA for compression or decompression, it may be not trivial to decide how many dimensions to keep. In this section we review a principled way to decide how many dimensions to keep.

The number of dimensions to keep is the only *hyper-parameter* of PCA. A method designed to decide how many dimensions/eigenvectors is the *proportion of variance*:

$$\text{POV} = \frac{\sum_{i=1}^m \lambda_i}{\sum_{j=1}^M \lambda_j},$$

where λ are eigenvalues, M is the dimensionality of the original data, and m is the chosen lower dimensionality.

Using the *POV* formula we may select a number M of dimensions/eigenvalues so that the proportion of variance is, for instance, equal to 95%.

Implement a new PCA for encoding and decoding that receives in input not the number of dimensions for projection, but the amount of proportion of variance to be preserved.

```
In [ ]: def encode_decode_pca_with_pov(A,p):
        # INPUT:
        # A      [NxM] numpy data matrix (N samples, M features)
        # p      float number between 0 and 1 denoting the POV to be preserved
        #
        # OUTPUT:
        # Ahat   [NxM] numpy PCA reconstructed data matrix (N samples, M features)
        # m      integer reporting the number of dimensions selected

        m = None
        Ahat = None

        return Ahat,m
```

Import the *lfw* dataset using the *get_lfw_data()* in *syntheticdata*. Use the implemented function to encode and decode the data by projecting on a lower dimensional space such that $\text{POV}=0.9$. Use the function *imshow* to plot and compare original and reconstructed pictures. Comment on the results.

```
In [ ]: X,y,h,w = syntheticdata.get_lfw_data()
```

```
In [ ]: Xhat,m = encode_decode_pca_with_pov(X,0.9)
```

```
In [ ]: plt.imshow(X[0,:].reshape((h, w)), cmap=plt.cm.gray)
plt.figure()
plt.imshow(Xhat[0,:].reshape((h, w)), cmap=plt.cm.gray)
```

Comment: Enter your comment here.

K-Means Clustering (Bachelor and master students)

In this section you will use the *k-means clustering* algorithm to perform unsupervised clustering. Then you will perform a qualitative assesment of the results.

Importing scikit-learn library

We start importing the module *cluster.KMeans* from the standard machine learning library *scikit-learn*.

```
In [353... from sklearn.cluster import KMeans
```

Loading the data

We will use once again the *iris* data set. The function *get_iris_data()* from the module *syntethicdata* returns the *iris* dataset. It returns a data matrix of dimension [150x4] and a label vector of dimension [150].

```
In [354... X,y = syntheticdata.get_iris_data()
```

Projecting the data using PCA

To allow for visualization, we project our data in two dimensions as we did previously. This step is not necessary, and we may want to try to use *k-means* later without the PCA pre-processing. However, we use PCA, as this will allow for an easy visualization.

```
In [355... _,P = pca(X,2)
```

Running k-means

We will now consider the *iris* data set as an unlabeled set, and perform clustering to this unlabeled set. We can compare the results of the clustering to the labeled classes.

Use the class *KMeans* to fit and predict the output of the *k-means* algorithm on the projected data. Run the algorithm using the following values of $k = \{2, 3, 4, 5\}$.

```
In [356... KM = KMeans(2)
yhat2 = KM.fit_predict(P)

KM = KMeans(3)
yhat3 = KM.fit_predict(P)

KM = KMeans(4)
yhat4 = KM.fit_predict(P)

KM = KMeans(5)
yhat5 = KM.fit_predict(P)
```

Qualitative assessment

Plot the results of running the k-means algorithm, compare with the true labels, and comment.

```
In [357... plt.scatter(P[:,0],P[:,1],c=yhat2)
plt.title('k=2')
plt.figure()

# Repeat for k=3, 4 and 5:
plt.scatter(P[:,0],P[:,1],c=yhat3)
plt.title('k=3')
plt.figure()


plt.scatter(P[:,0],P[:,1],c=yhat4)
plt.title('k=4')
plt.figure()
```






```
plt.scatter(P[:,0],P[:,1],c=yhat5)
plt.title('k=5')
plt.figure()

plt.figure()
plt.scatter(P[:,0],P[:,1],c=y)
plt.title('Original data')
```


Out[357]: Text(0.5, 1.0, 'Original data')







<Figure size 432x288 with 0 Axes>



Comment: Each plot shows k clusters.

Quantitative Assessment of K-Means (Bachelor and master students)

We used k-means for clustering and we assessed the results qualitatively by visualizing them. However, we often want to be able to measure in a quantitative way how good the clustering was. To do this, we will use a classification task to evaluate numerically the goodness of the representation learned via k-means.

Reload the *iris* dataset. Import a standard `LogisticRegression` classifier from the module `sklearn.linear_model`. Use the k-means representations learned previously (`yhat2,...,yhat5`) and the true label to train the classifier. Evaluate your model on the training data (we do not have a test set, so this procedure will assess the model fit instead of generalization) using the `accuracy_score()` function from the `sklearn.metrics` module. Plot a graph showing how the accuracy score varies when changing the value of k. Comment on the results.

- Train a Logistic regression model using the first two dimensions of the PCA of the iris data set as input, and the true classes as targets.
- Report the model fit/accuracy on the training set.
- For each value of K:
 - One-Hot-Encode the classes output by the K-means algorithm.
 - Train a Logistic regression model on the K-means classes as input vs the real classes as targets.
 - Calculate model fit/accuracy vs. value of K.
- Plot your results in a graph and comment on the K-means fit.

```
In [358... from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import OneHotEncoder

# Iris dataset
X,y = syntheticdata.get_iris_data()

# pca with 2 dimensions
_,P = pca(X,2)

# Training the classifier
log_reg_cl = LogisticRegression()
log_reg_cl.fit(P, y)
pred = log_reg_cl.predict(P)
```

```

acc_train = accuracy_score(pred, y)
print("Accuracy - training set: {}".format(acc_train))
encoder = OneHotEncoder()

# K means
K = [2, 3, 4, 5]
# List of accuracy per k
acc_scores = []
for k in K:
    KM = KMeans(k)
    yhat = KM.fit_predict(P)
    yhat_reshape = yhat.reshape(-1,1)
    # One-hot-encoder for K
    enc_k = encoder.fit_transform(yhat_reshape)
    # Training logistic classifier on k and target value y
    log_reg_cl.fit(enc_k, y)
    # Computing prediction
    pred_enc_k = log_reg_cl.predict(enc_k)
    # Computing accuracy
    acc = accuracy_score(pred_enc_k, y)
    acc_scores.append(acc)
    print("Accuracy when k = {}: {}".format(k, acc))

# Plotting accuracy
plt.plot(K, acc_scores)
plt.xlabel("K-value")
plt.ylabel("Accuracy")
plt.figure()


```

```

Accuracy - training set: 0.9666666666666667
Accuracy when k = 2: 0.6666666666666666
Accuracy when k = 3: 0.8866666666666667
Accuracy when k = 4: 0.84
Accuracy when k = 5: 0.9066666666666666

```

Out[358]: <Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>

Comment: The accuracy is high when $k=3$ because the original data has 3 classes. When $k > 3$ the accuracy is decreasing, but after a point it's increasing again. When $k=5$ the accuracy is higher. When increasing K to 5 it's possible that it was able to capture some additional structure in the data that was not captured by the original 3 classes, resulting in a higher accuracy score.

Conclusions

In this notebook we studied **unsupervised learning** considering two important and representative algorithms: **PCA** and **k-means**.

First, we implemented the PCA algorithm step by step; we then run the algorithm on synthetic data in order to see its working and evaluate when it may make sense to use it and when not. We then considered two typical uses of PCA: for **visualization** on the *iris* dataset, and for **compression-decompression** on the *lfw* dataset.

We then moved to consider the k-means algorithm. In this case we used the implementation provided by *scikit-learn* and we applied it to another prototypical unsupervised learning problem: **clustering**; we used *k-means* to process the *iris* dataset and we evaluated the results visually.

In the final part, we considered two additional questions that may arise when using the above algorithms. For PCA, we considered the problem of **selection of hyper-parameters**, that is, how we can select the hyper-parameter of our algorithm in a reasonable fashion. For

k-means, we considered the problem of the **quantitative evaluation** of our results, that is, how can we measure the performance or usefulness of our algorithms.